

Genomic Repeat inference from Depth (GRiD): A Pipeline for Haplotype-Resolved Large VNTR Copy Number Estimation

Zachary Caterer^{1,2,3}, Meng Lin¹, Qiang An⁴, Maizy S. Brasher¹, Elbay Aliyev¹, Harriet Dashnow¹, Joanne B. Cole^{1,3}, Ethan Lange¹, Mariaelisa Graff⁴, Christy L. Avery⁴, Christopher R. Gignoux^{1,3}, and Maggie Stanislawski^{1,3}

¹ Department of Biomedical Informatics, University of Colorado Anschutz, Aurora, CO USA ² Department of Chemical and Biological Engineering, University of Colorado Boulder, Boulder, CO USA ³ Interdisciplinary Quantitative Biology PhD Program, Biofrontiers Institute, University of Colorado Boulder, Boulder, CO USA ⁴ Department of Epidemiology, University of North Carolina Chapel Hill, Chapel Hill, NC USA ^{*} Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Variable number tandem repeats (VNTRs) are genomic regions in which a sequence motif is repeated in tandem a variable number of times across individuals. Large VNTRs, repeat units spanning hundreds to thousands of base pairs, are a substantial source of structural variation in the human genome and have been linked to complex traits and diseases (Mukamel et al., 2021). Because repeat units in these VNTRs can be more than 35 times the length of the standard 150-bp short sequencing reads, large VNTRs cannot be genotyped by conventional variant-calling approaches. Furthermore, accurate copy number estimation from short-read whole-genome sequencing (WGS) remains challenging. A well-characterized example is the Kringle IV type-2 (KIV-2) VNTR within the LPA gene, which is the primary genetic determinant of lipoprotein(a) [Lp(a)] concentration (Mukamel et al., 2021; Utermann et al., 1987). Lp(a) is a causal and highly heritable ($h^2 > 70\%$) atherosclerotic cardiovascular disease (ASCVD) risk factor elevated in an estimated 1.5 billion people globally (Reyes-Soffer et al., 2022). Both Lp(a) concentrations and effect of LPA genetic variants vary substantially across ancestry groups, underscoring the importance of ancestry-aware genetic analyses and risk prediction (Kronenberg et al., 2022; Nordestgaard et al., 2010; Tsimikas, 2017). The KIV-2 repeat unit spans ~5.5 kb, with individuals carrying 1–40 copies per haplotype, and KIV-2 copy number is inversely correlated with Lp(a) concentration (Clarke et al., 2009; Kamstrup et al., 2009; Mukamel et al., 2021; Utermann et al., 1987). Accurate estimation of KIV-2 copy number therefore has important implications for genetic studies of Lp(a) and cardiovascular disease, including analyses of ancestry-related variation and genetic risk (Kronenberg et al., 2022; Mukamel et al., 2021). While this VNTR serves as a region of great interest, prior work has been focused on custom pipelines that are difficult to apply in modern environments, limiting applications by the broader field.

GRiD (Genomic Repeat inference from Depth) is an open-source Python pipeline that addresses this methodological gap. GRiD estimates diploid and haplotype-resolved VNTR copy number from short-read WGS data. GRiD combines read-depth normalization, read depth profile-matched nearest-neighbor estimation, and identity-by-descent (IBD)-based haplotype inference (Hujoel et al., 2026). Although GRiD was developed and tested on the LPA KIV-2 locus, its core framework is locus-agnostic and applicable to any large VNTR with sufficient sample size. The pipeline is designed to be accessible to researchers without extensive bioinformatics

44 expertise.

45 Statement of need

46 Large VNTRs present a shared set of computational challenges that place them outside the
 47 scope of existing repeat genotyping software and standard genomic analysis pipelines. Their
 48 repeat units are orders of magnitude larger than the short tandem repeats (STRs) targeted
 49 by most available tools (Bakhtiari et al., 2018; Dolzhenko & others, 2019; Mousavi et al.,
 50 2019; Willems et al., 2017), and in individuals with high-copy-number alleles the VNTR
 51 region can span hundreds of kilobases, making read-pair spanning approaches difficult with
 52 standard paired-end reads. However, haplotype resolution of VNTR from short-read data is
 53 still important for downstream genomic analyses that depend on allele-specific copy number,
 54 including analyses of statistical associations, genomic coverage, and haplotype based tract,
 55 particularly given the expense and difficulty of obtaining long-read data for large epidemiological
 56 cohorts. These challenges can be directly addressed through three features: (1) a read-depth
 57 normalization and nearest-neighbor framework that is locus-agnostic and parameterized entirely
 58 through a configuration file, (2) an IBD-based iterative phasing algorithm (Hujoel et al., 2026)
 59 that decomposes diploid copy number into haplotype-specific estimates using cohort relatedness,
 60 and (3) a single installable pipeline that encapsulates the full workflow from short-read WGS
 61 to haplotype-resolved copy number in an intuitive, reproducible, multi-threaded form.

62 Although several software packages address aspects of large VNTRs analysis (Hujoel et al.,
 63 2026; Mukamel et al., 2021), no existing open-source tool provides an end-to-end workflow for
 64 estimating both diploid and haplotype-resolved copy number from short-read whole-genome
 65 sequencing data at modern biobank scales. Consequently, many studies have relied on custom
 66 analysis pipelines, limiting reproducibility, portability, and cross-cohort compatibility. GRiD
 67 fills this gap by providing a configurable, documented, and openly available implementation of
 68 these methods.

69 State of the field

70 Several mature software packages exist for genotyping tandem repeats from short-read
 71 sequencing data. ExpansionHunter (Dolzhenko & others, 2019), and HipSTR (Willems
 72 et al., 2017) were developed primarily for short tandem repeats (STRs) whereas GangSTR
 73 (Mousavi et al., 2019) was developed for short VNTRs (defined by a repeat motif length of up
 74 to 20bp). However, these methods are designed for repeat lengths where short reads can still
 75 provide sufficient direct evidence about repeat structure. Likewise, TRTools (Mousavi et al.,
 76 2021) provides downstream analysis utilities for STR genotypes but does not extend to large
 77 VNTR copy number estimation. adVNTR (Bakhtiari et al., 2018) supports genome-wide VNTR
 78 genotyping using hidden Markov models but is designed for repeat units that are substantially
 79 shorter (<150 bp) than loci such as the ~5.5 kb LPA KIV-2 repeat. Earlier methods such
 80 as VNTRseek (Gelfand et al., 2014) detected shorter VNTRs by comparing read-derived
 81 repeat counts to a reference catalog, while GtTR (Ganesamoorthy et al., 2018) estimated
 82 absolute target VNTR copy number by scaling short-read depth across locus boundaries against
 83 a baseline genotype derived from long reads, though it remains unable to phase individual
 84 diploid alleles. More recently, danbing-tk (Lu et al., 2021) introduced repeat-pangenome
 85 graphs for genome-wide profiling of VNTR length and motif composition from short-read
 86 sequencing data. This approach represents VNTR sequence diversity using haplotype-resolved
 87 assemblies and estimates VNTR dosage from graph-specific k-mer counts. Recent advances in
 88 long-read sequencing have enabled direct characterization of complex repeat regions (Showpnil,
 89 2024). For example, Kivvi (PacificBiosciences, 2026) estimates KIV-2 copy number from
 90 PacBio HiFi sequencing data. However, most large population cohorts (eg, 1000G (The
 91 1000 Genomes Project Consortium, 2015), TOPMed (Taliun et al., 2021), and the Genotype-
 92 Tissue Expression (GTEx) project (GTEx Consortium, 2017)) remain based on short-read

whole-genome sequencing, motivating methods that can estimate VNTR copy number without requiring long-read data. KILDA (Molitor & others, 2025) estimates KIV-2 diploid copy number directly from short-read FASTQ files using an alignment-free k-mer-counting approach normalized against non-repetitive sequences. The proprietary Illumina DRAGEN LPA Caller similarly estimates KIV-2 copy number but requires the FPGA-accelerated DRAGEN platform. vntr-calling-nf (genepi, 2024) provides a Nextflow workflow for VNTR analysis but focuses on coding variant detection rather than large-VNTR copy number quantification and has primarily been validated using whole-exome sequencing data. Collectively, these tools do not provide an end-to-end solution for estimating haplotype-resolved copy number of large VNTRs from short-read whole-genome sequencing data at population scale.

Software design

GRiD integrates approaches from prior work for estimating large VNTR copy number from short-read sequencing (Hujoel et al., 2026; Mukamel et al., 2021). It provides a configurable, end-to-end workflow for estimating both diploid and haplotype-resolved large VNTR copy numbers from short-read whole-genome sequencing data at cohort scale. The pipeline is designed to be accessible to researchers without extensive bioinformatics expertise.

GRiD is implemented in Python (≥ 3.8) and follows a modular architecture organized into seven sequential pipeline steps, each implemented as an independent utility module under `grid/utis/`. The pipeline is orchestrated by `grid/pipeline.py` and driven by a user-supplied YAML configuration file, which specifies file paths, genomic coordinates, per-step parameters, and which steps to execute. The command-line interface is implemented with Click (Ronacher & others, 2024) and provides an entry point: `grid wgs` for whole-genome sequencing data.

The Whole Genome Sequencing pipeline proceeds as follows:

1. **Index verification / creation:** Ensures CRAM/BAM index files (`.crai/.bai`) exist for all input samples, creating them with `samtools` (Li et al., 2009) if needed.
2. **Read counting:** Counts properly paired reads in the target VNTR region using `pysam` (Heger & others, 2024), filtering by mapping quality and SAM flags.
3. **Coverage estimation:** Runs `mosdepth` (Pedersen & Quinlan, 2018) across the genome in binned mode to produce per-sample depth profiles.
4. **Coverage normalization:** Normalizes coverage within individuals (by sample mean depth) and across individuals (by a z-score transformation), then filters to high-variance regions enriched for VNTR signal.
5. **Nearest-neighbor identification:** Uses `scikit-learn` (Pedregosa et al., 2011) to compute Euclidean distances in normalized depth space and identify the top-N genomically similar neighbors per individual, providing a cohort-matched reference for CNV normalization.
6. **Diploid copy number estimation:** Computes per-exon diploid copy number for each individual by normalizing their read counts relative to their nearest neighbors' read counts, accounting for sample-specific sequencing depth.
7. **Haplotype inference:** Applies an IBD-based iterative phasing algorithm (Hujoel et al., 2026) that uses haplotype-matched neighbors to decompose diploid copy number estimates into haplotype-specific contributions.

All multi-sample steps are parallelized using Python's `concurrent.futures.ThreadPoolExecutor`, and progress is reported via the Rich terminal library (McGugan, 2024). Configuration validation at pipeline startup provides early, informative error messages before any compute-intensive steps run.

Research impact statement

GRiD has been applied to >50,000 study participants with whole-genome sequencing data from multiple TOPMed cohorts (Taliun et al., 2021) and 1000 Genomes Project (The 1000 Genomes Project Consortium, 2015), comprising participants across continental populations of African, European, and American ancestry (including self-reported Hispanic/Latino individuals and individuals with indigenous American genetic ancestry). These analyses form the basis of ongoing research manuscripts and demonstrate the applicability of GRiD to population-scale analyses and heterogeneous sequencing datasets.

GRiD provides a reproducible, open-source implementation of a previously custom analysis workflow, enabling large-VNTR copy number estimation from existing short-read whole-genome sequencing datasets. By packaging the complete workflow with documentation, example datasets, and automated installation, GRiD reduces the technical barriers to incorporating large VNTR analyses into population-scale genomic studies.

Mathematics

The mathematical framework implemented in GRiD follows previously published methods. Coverage estimation, normalization, and diploid copy-number estimation follow Mukamel et al. (Mukamel et al., 2021), while haplotype-specific copy-number inference follows Hujoel et al. (Hujoel et al., 2026). The equations below describe their implementation within GRiD; detailed methodological derivations and justification are provided in the original publications.

Coverage estimation. Following Mukamel et al. (Mukamel et al., 2021, Supplemental Materials, Section 2), weighted mean depth across the VNTR region from mosdepth (Pedersen & Quinlan, 2018) bins is computed as:

$$\text{Coverage} = \left\lceil 100 \cdot \frac{\sum_{i \in R} (\bar{C}_i \cdot (\min(e, r_{i,e}) - \max(s, r_{i,s})))}{\sum_{i \in R} (\min(e, r_{i,e}) - \max(s, r_{i,s}))} \right\rceil$$

where R is the set of bins overlapping the region $[s, e]$, $\bar{C}_i$ is the mean depth of bin i , and $r_{i,s}$ and $r_{i,e}$ are the bin boundaries.

Normalization. Following Mukamel et al. (Mukamel et al., 2021, Supplemental Materials, Section 2), let $D \in \mathbb{R}^{N \times M}$ be the depth matrix for N individuals across M genomic bins. Within-individual normalization gives $D_{ij}^{(1)} = D_{ij} / \bar{D}_i$, where $\bar{D}_i$ is the individual mean. Across-individual normalization then yields:

$$D_{ij}^{\text{norm}} = \frac{D_{ij}^{(1)} - \mu_j}{\sqrt{\mu_j}}$$

where μ_j is the population mean of $D_{ij}^{(1)}$. Regions are filtered by a variance ratio σ_j^2 / μ_j , and the top fraction of high-variance regions are retained for neighbor computation.

Neighbor-normalized diploid copy number. GRiD uses the coverage-profile normalization approach described by Mukamel et al. (Mukamel et al., 2021, Supplementary Text, Section 2). For each individual, VNTR read depth is normalized relative to a set of samples with similar genome-wide coverage profiles. GRiD identifies these reference samples using nearest-neighbor search in the normalized depth space, implemented with scikit-learn (Pedregosa et al., 2011). The resulting diploid copy-number estimate for individual i is:

$$\widehat{\text{dipCN}}_i = \frac{r_i/s_i}{\frac{1}{|N_i|} \sum_{j \in N_i} r_j/s_j}$$

where N_i is the set of nearest-neighbor samples for individual i , and r_i/s_i represents the individual's VNTR read count normalized by sequencing depth.

Haplotype inference. Haplotype-specific copy-number inference follows the iterative phasing approach described by Hujoel et al. (Hujoel et al., 2026, Supplemental Materials, Section 5.3). In this framework, diploid VNTR copy-number estimates are partitioned between an individual's two haplotypes using information from IBD-matched haplotypes, with haplotype-specific estimates updated iteratively. GRiD formalizes this procedure as:

$$h_{i,k}^{(t+1)} = \widehat{\text{dipCN}}_i \cdot \frac{\bar{h}_{N_i,k}^{(t)}}{\bar{h}_{N_i,1}^{(t)} + \bar{h}_{N_i,2}^{(t)}}$$

where $h_{i,k}^{(t)}$ is the estimated copy number of haplotype k for individual i at iteration t , and $\bar{h}_{N_i,k}^{(t)}$ is the mean haplotype copy number among IBD-matched reference haplotypes on haplotype k . At each iteration, the relative copy numbers of the IBD-matched haplotypes determine how the individual's diploid estimate is allocated between its two haplotypes. The estimates are then updated iteratively until convergence.

AI Usage Disclosure

The authors used GitHub Copilot and Anthropic Claude (Claude Sonnet 4) during the development of this software and the preparation of this manuscript. GitHub Copilot was used as an interactive coding assistant to suggest code completions, refactoring ideas, and boilerplate implementations. Claude was used to assist with code review, debugging, documentation development, and editorial improvements to the manuscript, including suggestions for clarity, organization, and grammar. AI tools were not used to generate scientific conclusions or interpret results. All AI-assisted code and manuscript edits were reviewed, validated, and, where necessary, modified by the authors before inclusion. The authors take full responsibility for the accuracy, correctness, and content of the software and manuscript.

Acknowledgements

We are deeply indebted to Drs. Po-Ru Loh and Margaux Hujoel for sharing code and providing additional technical support. We would additionally like to thank Dr. Akshay Avvaru for his help and support, and Dr. Daniel Larremore for his careful review and for catching several typos before they could escape into the wild. This work was supported by NIH K01: HL157658, R01: HL172887 and NSF: 2022138 grants. The funders had no role in the design, development, analysis, or preparation of this work.

References

- Bakhtiari, M., Shleizer-Burko, S., Gymrek, M., Bansal, V., & Bafna, V. (2018). Targeted genotyping of variable number tandem repeats with adVNTR. *Genome Research*, 28(11), 1709–1719. <https://doi.org/10.1101/gr.235119.118>
- Clarke, R., Peden, J. F., Hopewell, J. C., & others. (2009). Genetic variants associated with Lp(a) lipoprotein level and coronary disease. *New England Journal of Medicine*, 361(26), 2518–2528. <https://doi.org/10.1056/NEJMoa0902604>

- 210 Dolzhenko, E., & others. (2019). ExpansionHunter: A sequence-graph-based tool to analyze
211 variation in short tandem repeat regions. *Bioinformatics*, 35(22), 4754–4756. <https://doi.org/10.1093/bioinformatics/btz431>
212
- 213 Ganesamoorthy, D., Cao, M. D., Duarte, T. P. S., Chen, W., & Coin, L. J. M. (2018). GtTR:
214 Bayesian estimation of absolute tandem repeat copy number using sequence capture and high
215 throughput sequencing. *BMC Bioinformatics*, 19(1), 267. [https://doi.org/10.1186/s12859-](https://doi.org/10.1186/s12859-018-2282-3)
216 [018-2282-3](https://doi.org/10.1186/s12859-018-2282-3)
- 217 Gelfand, Y., Hernandez, Y., Loving, J., & Benson, G. (2014). VNTRseek—a computational
218 tool to detect tandem repeat variants in high-throughput sequencing data. *Nucleic Acids*
219 *Research*, 42(14), 8884–8894. <https://doi.org/10.1093/nar/gku642>
- 220 genepi. (2024). *Vntr-calling-nf: A Nextflow pipeline for VNTR variant calling*. <https://github.com/genepi/vntr-calling-nf>
221
- 222 GTEx Consortium. (2017). Genetic effects on gene expression across human tissues. *Nature*,
223 550(7675), 204–213. <https://doi.org/10.1038/nature24277>
- 224 Heger, A., & others. (2024). *Pysam: A python module for reading and manipulating*
225 *SAM/BAM/CRAM files*. <https://github.com/pysam-developers/pysam>
- 226 Hujoel, M. L., Handsaker, R. E., Tang, D., Kamitaki, N., Mukamel, R. E., Rubinacci, S.,
227 Palamara, P. F., McCarroll, S. A., & Loh, P.-R. (2026). Insights into DNA repeat expansions
228 among 900,000 biobank participants. *Nature*, 1–10. [https://doi.org/10.1038/s41586-025-](https://doi.org/10.1038/s41586-025-09886-z)
229 [09886-z](https://doi.org/10.1038/s41586-025-09886-z)
- 230 Kamstrup, P. R., Tybjaerg-Hansen, A., Steffensen, R., & Nordestgaard, B. G. (2009). Genetically
231 elevated lipoprotein(a) and increased risk of myocardial infarction. *JAMA*, 301(22),
232 2331–2339. <https://doi.org/10.1001/jama.2009.801>
- 233 Kronenberg, F., Mora, S., Stroes, E. S. G., & others. (2022). Lipoprotein(a) in atherosclerotic
234 cardiovascular disease and aortic stenosis: A European Atherosclerosis Society consensus
235 statement. *European Heart Journal*, 43(39), 3925–3946. [https://doi.org/10.1093/](https://doi.org/10.1093/eurheartj/ehac361)
236 [eurheartj/ehac361](https://doi.org/10.1093/eurheartj/ehac361)
- 237 Li, H., Handsaker, B., Wysoker, A., & others. (2009). The sequence alignment/map format and
238 SAMtools. *Bioinformatics*, 25(16), 2078–2079. [https://doi.org/10.1093/bioinformatics/](https://doi.org/10.1093/bioinformatics/btp352)
239 [btp352](https://doi.org/10.1093/bioinformatics/btp352)
- 240 Lu, J.-H., Yang, C.-C., Chou, E. J., Gerasimova, A., Melendez, J., Markello, A., & Chaisson,
241 M. J. P. (2021). Profiling variable-number tandem repeat variation across populations
242 using repeat-pangenome graphs. *Nature Communications*, 12(1), 4250. [https://doi.org/](https://doi.org/10.1038/s41467-021-24378-0)
243 [10.1038/s41467-021-24378-0](https://doi.org/10.1038/s41467-021-24378-0)
- 244 McGugan, W. (2024). *Rich: Rich text and beautiful formatting in the terminal*. <https://github.com/Textualize/rich>
245
- 246 Molitor, L., & others. (2025). KILDA: K-mer-based inference of LPA diploid copy number
247 from short-read alignment-free data. *NAR Genomics and Bioinformatics*. [https://doi.org/](https://doi.org/10.1093/nargab/lqaf070)
248 [10.1093/nargab/lqaf070](https://doi.org/10.1093/nargab/lqaf070)
- 249 Mousavi, N., Margoliash, J., Pusarla, N., & others. (2021). TRTools: A toolkit for genome-
250 wide analysis of tandem repeats. *Bioinformatics*, 37(18), 3062–3064. [https://doi.org/10.](https://doi.org/10.1093/bioinformatics/btaa736)
251 [1093/bioinformatics/btaa736](https://doi.org/10.1093/bioinformatics/btaa736)
- 252 Mousavi, N., Shleizer-Burko, S., Yanicky, R., & Gymrek, M. (2019). GangSTR: Template-
253 based estimation of short tandem repeats. *Nucleic Acids Research*, 47(15), e90. <https://doi.org/10.1093/nar/gkz501>
254
- 255 Mukamel, R. E., Handsaker, R. E., Sherman, M. A., & others. (2021). Protein-coding repeat
256 polymorphisms strongly shape diverse human phenotypes. *Science*, 373(6562), 1499–1505.

- 257 <https://doi.org/10.1126/science.abg8289>
- 258 Nordestgaard, B. G., Chapman, M. J., Ray, K., Borén, J., Andreotti, F., Watts, G. F.,
259 Ginsberg, H. N., Amarengo, P., Catapano, A., Descamps, O. S., Fisher, E., Kovanen, P. T.,
260 Kuivenhoven, J. A., Lesnik, P., Masana, L., Reiner, Z., Taskinen, M.-R., Tokgözoğlu, L., &
261 Tybjærg-Hansen, A. (2010). Lipoprotein(a) as a cardiovascular risk factor: Current status.
262 *European Heart Journal*, 31(23), 2844–2853. <https://doi.org/10.1093/eurheartj/ehq386>
- 263 PacificBiosciences. (2026). Kivvi: KIV-2 and D4Z4 copy number estimation from PacBio HiFi
264 whole-genome sequencing. <https://github.com/PacificBiosciences/Kivvi>
- 265 Pedersen, B. S., & Quinlan, A. R. (2018). Mosdepth: Quick coverage calculation for genomes
266 and exomes. *Bioinformatics*, 34(5), 867–868. <https://doi.org/10.1093/bioinformatics/btx699>
- 267
- 268 Pedregosa, F., Varoquaux, G., Gramfort, A., & others. (2011). Scikit-learn: Machine learning
269 in Python. *Journal of Machine Learning Research*, 12, 2825–2830. <https://jmlr.org/papers/v12/pedregosa11a.html>
- 270
- 271 Reyes-Soffer, G., Ginsberg, H. N., Berglund, L., Duell, P. B., Heffron, S. P., Kamstrup,
272 P. R., Lloyd-Jones, D. M., Marcovina, S. M., Yeang, C., Koschinsky, M. L., & others.
273 (2022). Lipoprotein (a): A genetically determined, causal, and prevalent risk factor for
274 atherosclerotic cardiovascular disease: A scientific statement from the american heart
275 association. *Arteriosclerosis, Thrombosis, and Vascular Biology*, 42(1), e48–e60. <https://doi.org/10.1161/ATV.0000000000000147>
- 276
- 277 Ronacher, A., & others. (2024). Click: Python composable command line interface toolkit.
278 <https://palletsprojects.com/p/click/>
- 279 Showpnil, E. H. G., I. A. (2024). Long-read sequencing resolves complex repeat structures in
280 the human genome. *Nature Genetics*. <https://doi.org/10.1038/s41525-024-00454-4>
- 281 Taliun, D., Harris, D. N., Kessler, M. D., Carlson, J., Szpiech, Z. A., Torres, R., Taliun, S. A.
282 G., Corvelo, A., Gogarten, S. M., Kang, H. M., Pitsillides, A. N., LeFaive, J., Lee, S.-B.,
283 Tian, X., Browning, B. L., Das, S., Emde, A.-K., Clarke, W. E., NHLBI Trans-Omics
284 for Precision Medicine (TOPMed) Consortium, & Abecasis, G. R. (2021). Sequencing of
285 53,831 diverse genomes from the NHLBI TOPMed Program. *Nature*, 590(7845), 290–299.
286 <https://doi.org/10.1038/s41586-021-03205-y>
- 287 The 1000 Genomes Project Consortium. (2015). A global reference for human genetic variation.
288 *Nature*, 526(7571), 68–74. <https://doi.org/10.1038/nature15393>
- 289 Tsimikas, S. (2017). A test in context: Lipoprotein(a): Diagnosis, prognosis, controversies,
290 and emerging therapies. *Journal of the American College of Cardiology*, 69(6), 692–711.
291 <https://doi.org/10.1016/j.jacc.2016.11.042>
- 292 Utermann, G., Menzel, H. J., Kraft, H. G., Duba, H. C., Kemmler, H. G., & Seitz, C. (1987).
293 Lp(a) glycoprotein phenotypes: Inheritance and relation to Lp(a)-lipoprotein concentrations
294 in plasma. *Journal of Clinical Investigation*, 80(2), 458–465. <https://doi.org/10.1172/JCI113093>
- 295
- 296 Willems, T., Zuk, O., Erlich, Y., & Gymrek, M. (2017). Genome-wide profiling of heritable
297 and de novo STR variations. *Nature Methods*, 14, 590–592. <https://doi.org/10.1038/nmeth.4267>
- 298